RL in NL

A little reinforcement learning example in NetLogo

Nicolas Payette

SCHOOL OF GEOGRAPHY
AND THE ENVIRONMENT

UNIVERSITY OF
OXFORD

June 4th 2020

Use reinforcement learning when...

Agents' **actions** generate **rewards** from the **environment** over time.

It can handle delayed reward: when actions taken **now** can affect **future** rewards.

But it's trial-and-error search: agents have to try different actions to evaluate their effects.

Agent-environment interaction in RL

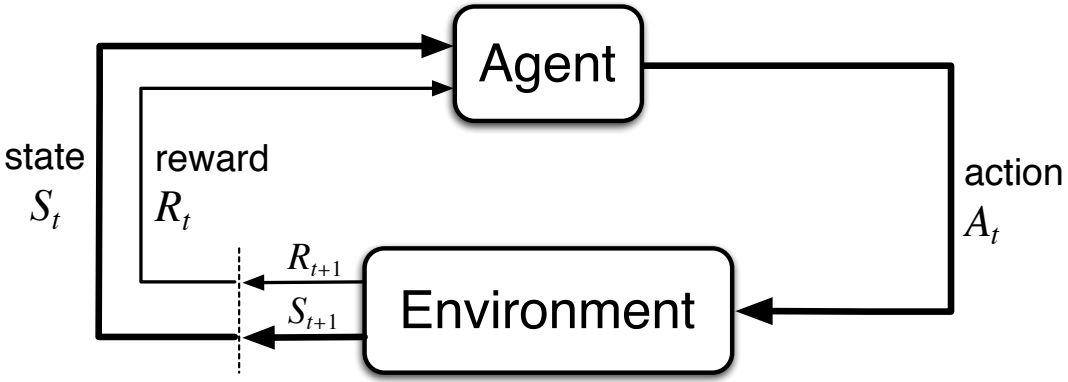

Sutton & Barto (2017) *Reinforcement Learning: An Introduction*

Temporal difference learning: TD(0)

Actions are chosen using an ϵ-**greedy** policy:

$$A \leftarrow \begin{cases} \text{a random action}, & \text{with probability } \epsilon \\ \text{the action that maximizes } V(S'), & \text{otherwise} \end{cases}$$

Agents must learn the value of the possible **states**:

$$V(S) \leftarrow V(S) + \alpha\Big(\overbrace{R + \gamma V(S') - V(S)}^{\text{TD error}} \Big)$$

Temporal difference learning: TD(0)

Actions are chosen using an ϵ**-greedy** policy:

$$A \leftarrow \begin{cases} \text{a random action}, & \text{with probability } \epsilon \\ \text{the action that maximizes } V(S'), & \text{otherwise} \end{cases}$$

Agents must learn the value of the possible **states**:

$$V(S) \leftarrow V(S) + \alpha\Big(\overbrace{R + \gamma V(S') - V(S)}^{\text{TD error}} \Big)$$

Constraint

Only works if the agent has a model of the environment: if it knows which action will bring about which state.

The Q function

$$Q : State \times Action \rightarrow Value$$

SARSA (on-policy – when you care about reward *while* learning):

$$Q(S,A) \leftarrow \underbrace{Q(S,A)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \Big(\overbrace{\underbrace{R}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{Q(S',A')}_{\textbf{value of next action taken}} - Q(S,A)}^{\text{TD error}} \Big)$$

Q-Learning (off-policy – when you just care about finding the optimal policy):

$$Q(S,A) \leftarrow \underbrace{Q(S,A)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \Big(\overbrace{\underbrace{R}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(S',a)}_{\textbf{value of best possible action}} - Q(S,A)}^{\text{TD error}} \Big)$$

Still ϵ-greedy, but this time we choose the action with the best Q-value.

Example: the windy grid world

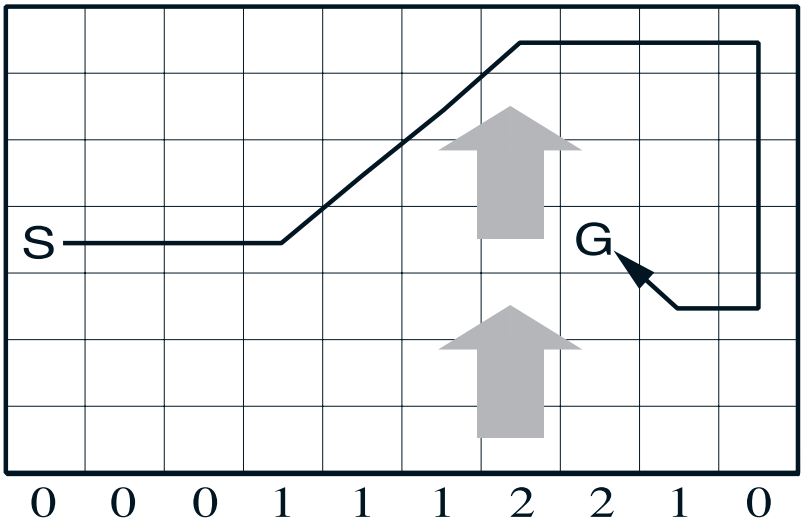

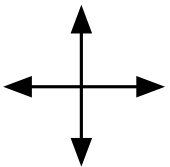

Actions

Sounds like a good candidate for a NetLogo model?